

Artificial Intelligence in a Human World

The AI's job is *plausibility*, not truth.

The AI said it compiles, the compiler says no.

Act 1

Rules ensure intent

Rules must be *truthful to the intent* and *applied consistently*.

“Rewrite this email in a friendlier tone without changing the meaning.”



Tools

Fast ▾



Rules must be *truthful* to the intent and applied consistently.

Visible Ink: What – the dialogue



*“Rewrite **this email** in a friendlier tone without changing the meaning.”*



Tools

Fast ▾



Rules must be *truthful to the intent* and *applied consistently*.

Visible Ink: *What* – the dialogue

The Story: *Why* – the goal

“Rewrite *this email* in a *friendlier tone* without changing the meaning.”



Tools

Fast ▾



Rules must be *truthful to the intent* and *applied consistently*.

Visible Ink: *What* – the dialogue

The Story: *Why* – the goal

“Rewrite *this email* in a *friendlier tone* without
changing the meaning.”

+ Tools

Fast ▾



Invisible Ink: *How* – rules & meaning

Rules must be *truthful to the intent* and *applied consistently*.

Visible Ink: *What* – the dialogue

The Story: *Why* – the goal

"Rewrite this email in a friendlier tone without changing the meaning."

+ Tools

Fast ▾



Invisible Ink: *"Don't – be buggy."*

Invisible Ink: *How* – rules & meaning

Hallucination is the feature, not a bug.

LLMs are *probabilistic engines* that are optimized to say
what sounds right rather than *what is right*.

– And we *want that!*

Hallucination is the **feature**, not a bug.

*(Unless it is **factually wrong**... then it's **a bug**, aka **a smudge**)*

LLMs are **probabilistic engines** that are **optimized** to say what **sounds right** rather than **what is right**.

– And we *want that!*

Rules need to be *observationally true* to the story.

User Journeys are Stories: (*Visible Ink*)

- Sign Up & Check Out
- Quarterly Report
- Automating Customer Service
- Analyzing large datasets to find trends

Stories have rules: (*Invisible Ink*)

- ● Minimize user effort & Build trust
- ● Requirements from regulatory bodies
- ● A Seamless, Omnichannel Experience
- ● patterns to benchmarks

Epistemics – what the model can say that is
storybound and context-limited.

Governance – what the system can trust/act
on with *deterministic, judge and receipts*.

Rules need to be *observationally true* to the story.

Rules *do not* need to be *objectively true*.

User Journeys are Stories: (*Visible Ink*)

- A Princess's journey to independence
- Protects humanity and fight for justice
- Hope comes from the most unlikely places
- Terminator; from destroyer to protector

Stories have rules: (*Invisible Ink*)

- In Disney – Animals can talk
- In DC – Superman can fly
- The one ring is evil and hobbits are pure
- What if a gun didn't want to be a gun?

Epistemics – what the model can say that is
storybound and context-limited.

Governance – what the system can trust/act
on with *deterministic, judge and receipts*.

If a superman movie turned into a batman movie, it wouldn't be a very good superman movie...



Tools

Fast ▾



“Rewrite this email in a friendlier tone without changing the meaning.”



Tools

Fast ▾



If a superman movie turned into a batman movie, it wouldn't be a very good superman movie...

+ Tools

The Feature: The Story/Goal

The Translation: The AI Alignment

Fast



"Rewrite this email in a friendlier tone without changing the meaning."

+ Tools

Fast



The Translation: "Please don't be buggy."

The Hallucination Paradox

The Paradox: Because probability is a *failure of governance*, you shouldn't try to "fix" the model to stop hallucinations; instead, *you must fix the boundary around it*.

The Feature

- *Probability is the engine of utility*
- When *content is aligned* to the goal
- *Plausible, story-bound* content
- *High Quality / Useful Content*

The Bug

- *Probability is a failure of governance* .
- When *content is misaligned* to the goal
- *Factually inaccurate* content
- *Low Quality / AI Smudge*

- **Visible Ink**: “*The Surface*”: *What the AI says*. It is the **dialogue**, the **tone**, and the **fluency**.
- **Invisible Ink**: “*The Understructure or armatage*”: *The rules the AI must obey*. It is the **business logic**, the **regulatory rules**, and your **specific intent**.
- **A Story**: Any process where the "Invisible Ink" (the story rules) **must govern** the "Visible Ink" (the output) to maintain trust and utility.
- **A Smudge**: The “*The Hallucination Paradox*” or "**when is a bug not a bug?**" – When the *Visible Ink* (the *plausible* performance) bleeds through and *overwrites the Invisible Ink* (the story/business rules).

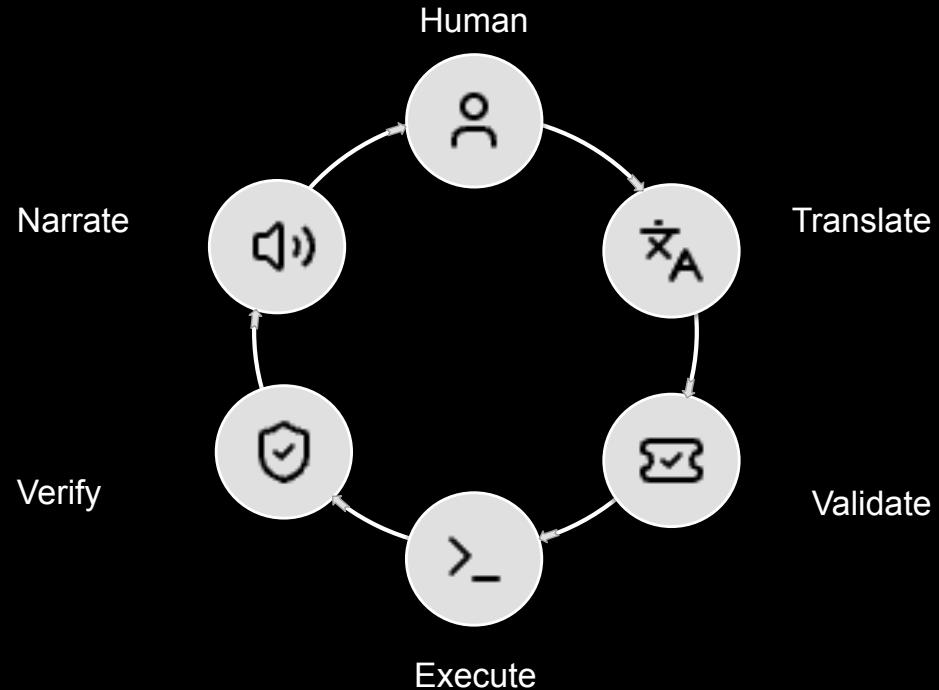
Act 2

Better boundaries

The Solver-Checker algorithm – align intent with rules

Goals – Protect Compliance and Narrative.

- ✓ Keep the Human in the Loop
- ✓ Align AI/Human Intent
- ✓ Keep AI in the middle
- ✓ Enable Agile AI



1. Grounding Pattern — "Shared Language" established before execution

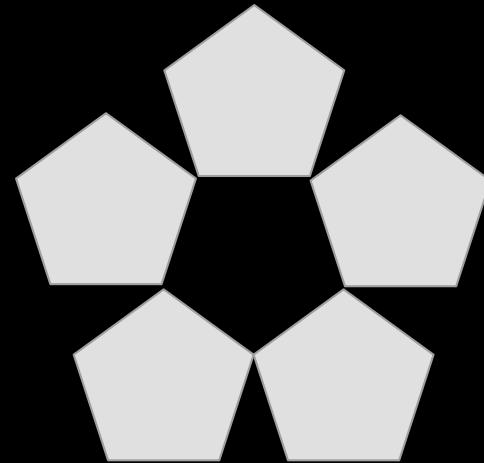
WHAT: Grounding is the anchor

WHEN: Policies, rules, lore; text-to-source-of-truth

WHY: Avoid invented facts and GIGO

HOW: Retrieve anchors + context bundle + execution
of tool + knowledge graphs + RAG + VectorDB

PROOF: Citations + “unknown” if missing



👍 Retrieve anchors + citations → say ‘unknown’ if missing

👎 Grounding failures are usually hidden, not obvious

2. Orchestration Pattern — coordinating steps in a controlled sequence

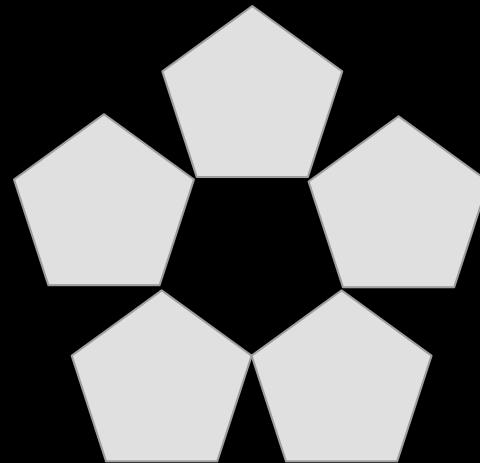
WHAT: Stepwise tool workflow

WHEN: Multi-step procedures

WHY: Prevent skipped/hidden actions

HOW: State machine + token/time budgets + rules

PROOF: "prompt engineering" to "systems engineering."



👍 It is testable, debuggable, and governable → example; accordion editing

👎 Orchestration chains add latency → keep steps minimal and set a latency budget

3. Verification Pattern — deterministic checks and proof of work

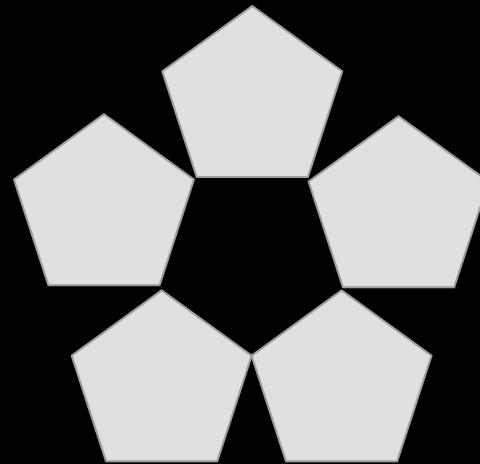
WHAT: Tests decide

WHEN: Binary correctness matters

WHY: Stop plausible wrong outputs

HOW: Deterministic judge boundary + Tools

PROOF: Pass/fail receipts



👍 Blind measure model performance objectively → If subjective, use a rubric or reference

👎 Trust vibes. Cherry-picked demos. No pass/fail receipts. Model decides without a judge

4. Trust UX Pattern — evidence and recovery options with user in control

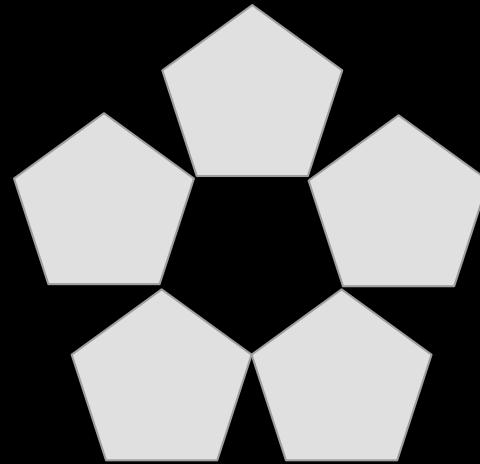
WHAT: Make uncertainty explicit

WHEN: Users: approve, override, correct

WHY: Avoid false confidence, (AI's dark triad)

HOW: Message + receipts; explicit fallbacks

PROOF: Scope, reasons, choices made, gaps shown



👍 Evidence-based review systems improve trust in decisions

👎 Jagged frontier; silent failures, false confidence, “no receipts, but looks right”

5. Compound Learning Pattern — small improvements over time

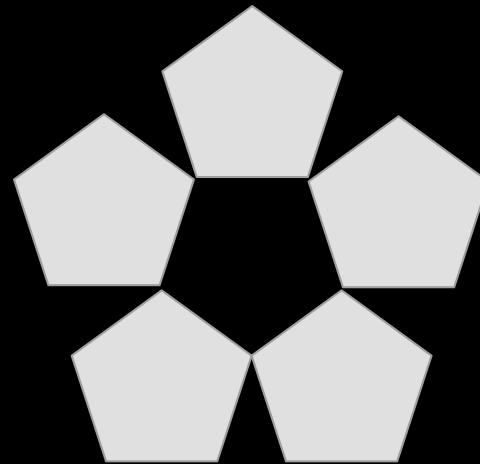
WHAT: Improve without regressions

WHEN: Generate can be scored

WHY: Prevent drift, scale performance

HOW: Offline eval harness + scorecards

PROOF: Score deltas, regression list, diary studies



- 👍 Self evaluation can drive rapid improvement → *This is far more important than most realize*
- 👎 Don't iterate by vibe; use fixed test sets

Solver-Checker Algorithm \leftrightarrow Core AI Design Patterns

Algorithmic Step: (what)	Design Pattern: (how)	Reason: (why)
Translate	1. Grounding	Protects Meaning
Validate	2. Orchestration	Protects Order
<i>Execute</i>	> Classical Software_	<i>Protects The Goal</i>
Verify	3. Verification	Protects Execution Truth
Narrate	4. Narrate	Protects Narrative Truth
<i>Outside</i> [#]	5. Learning	<i>Protects Progress</i>

👍 Testing in the Loop

- Deterministic unit tests – true unit tests
- Contract tests for the model – unit-test-like, but not text-equality
- Eval regression tests – the real safety net, proof-of-work
- *Don't unit test creativity – test tools, contracts, and regressions*

👍 AI Testing in the Loop - Practical rules

- Never let the model-judge be the only gate for correctness
- Calibrate the judges with “Golden Sets”
- Prefer pairwise ranking over absolute scoring
- Reduce correlated failure and find edge cases
- *A/B tests (Go talk to marketing!)*
- In-process testing > fire and then forget (unit tests)



AI Anti-Patterns (common pitfalls in AI development)

1. Single-Shots prompts *“One prompt, one hope.”*
2. One prompt, one judge *Grading your own papers*
3. The “god” prompt *Epistles and “thou shalt not...”*
4. Iteration by vibes *Cargo-cults – ritual inclusion that serves no purpose*
5. Waiting for AGI *AGI is asymptotic to perfect plausibility*

Act 3

Proof of work

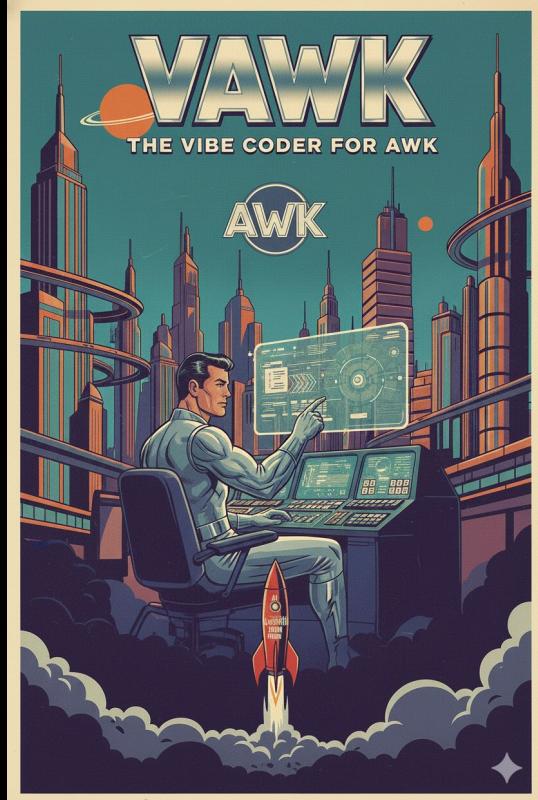
Demo 1: VAWK – AWK vibe coding

Governance – what the system can trust/act on with deterministic, judge and receipts.

1. **Grounding:** Backus-Naur Form (BNF) check
2. **Orchestration:** Propose → RAG → Run → Patch
3. **Verification:** Interpreter + tests decide
4. **Trust UX:** Receipts are visible
5. **Learning:** Regression sets



<https://github.com/dwellman/vawk>



Demo 2: A BUUI Adventure

Epistemics -- what the model can say that is storybound and context-limited. (Who is allowed to say 'this is correct.')

1. Grounding: World anchored in state transition

2. Orchestration: One command per tick

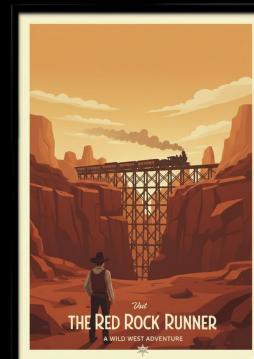
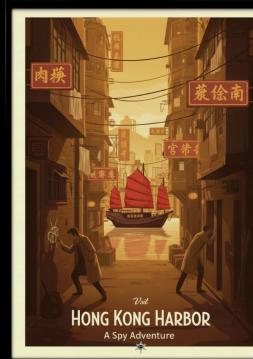
3. Verification: Game engine-rules decide via RAG

4. Trust UX: State change w/ receipts

5. Learning: Scenario replay with eval



<https://github.com/dwellman/adventure>

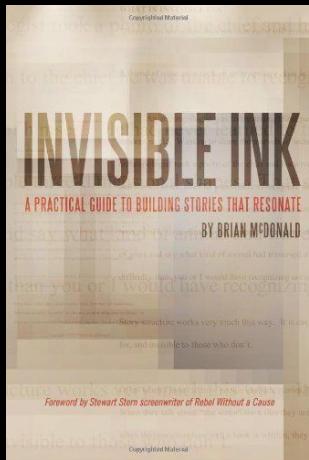


Demos: A Recap 1980s-era software *now with AI!*

Pattern	VAWK (Coding/Execution Truth)	BUUI (Gaming/Narrative Truth)
1. Grounding	The Syntax: BNF Grammars & Compiler Rules.	The World-State: Player inventory & Location DB.
2. Orchestration	The Loop: Solver proposes code; Checker tests it.	The DM: Narrative engine tracks "invisible" game rules.
3. Verification	The Judge: An external Python interpreter/compiler.	The Rulebook: RAG-lookup to ensure actions are "legal."
4. Narrative	The Clean-up: Translating raw logs into a "Success" message.	The Story: Turning state changes into immersive prose.
5. Receipts	The Log: A full trace of why the code was patched.	The Game State: A reflective history of every world-state change.

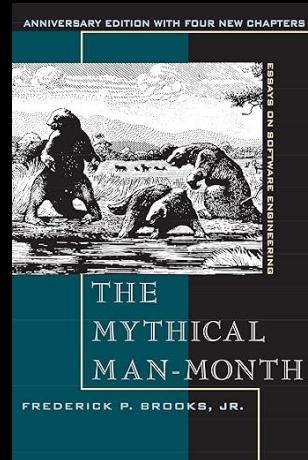
Invisible Ink:

- Brian McDonald



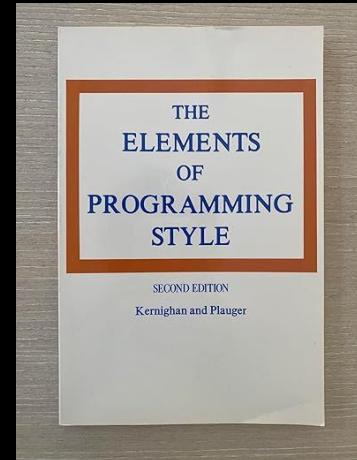
The Mythical Man-Month:

- Frederick Brooks Jr.



The Elements of Programming Style

- Brian W. Kernighan, P. J. Plauger



<https://a.co/d/8zI0QIJ>



<https://a.co/d/4Rc07fq>



<https://a.co/d/35IA5ZW>

Q&A

The AI's job is **plausibility**, not **truth**.

Building the boundary and injecting the facts. *That's your job.*

- 1. Grounding: AI-assisted triage can notify specialists from imaging workflows.**
 - AI “parallel stroke workflow” tool and workflow timing measures. [[AHA Journals](#)]
 - LVO detection software and time-to-treatment/outcomes. [[JAMA Network](#)]
- 2. Orchestration: AI-assisted stroke triage can notify specialists from imaging workflows.**
 - AI “parallel stroke workflow” tool and workflow timing measures. [[AHA Journals](#)]
 - LVO detection software and time-to-treatment/outcomes. [[JAMA Network](#)]
- 3. Verification: Standardized benchmarks**
 - HELM (multi-metric benchmarking and transparency).
 - BIG-bench (broad task suite; human baselines; calibration discussion). [[arXiv:2206.04615](#)]
- 4. Trust UX: Evidence-based review systems improve trust in decisions.**
 - Trust in automation review [[SAGE Journals](#)]
 - Algorithm aversion [[sol3:2466040](#)]
 - The Impact of Placebic Explanations [[eiband2019chiea](#)]
- 5. Learning: Self evaluation can drive rapid improvement.**
 - Self-Refine: Iterative Refinement with Self-Feedback [[arXiv:2303.17651](#))
 - Reflexion (self-reflection + memory improves agent performance) [[arXiv:2303.11366](#)]
 - Constitutional AI (self-critique/-revision framed as AI feedback during training). [[arXiv:2212.08](#)]

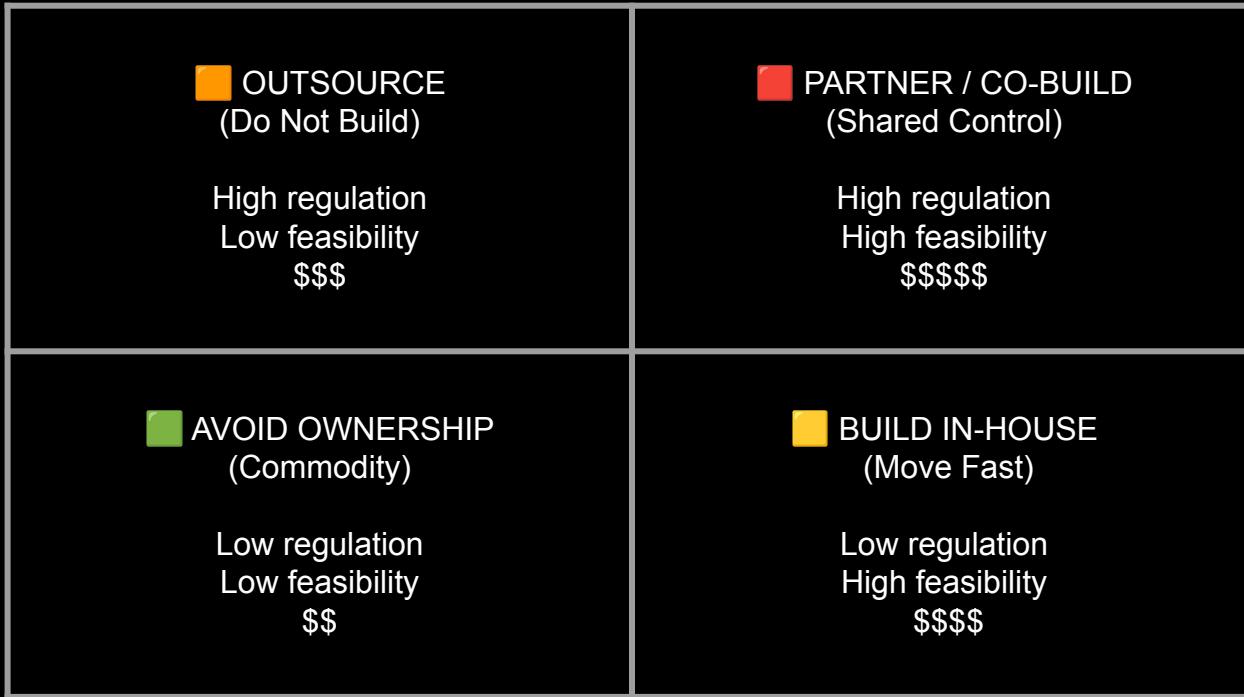
Presentation Review:

- **Thesis:** The AI's job is plausibility, not truth.
- **Keystone:** AI says it compiles. The compiler says no.
- **Patterns:** 1. Grounding, 2. Orchestration 3. Verification 4. Trust UX 5. Learning
- **Solver-Checker:** Translate → Validate → Execute → Verify → Narrate

Position Papers:

- **Move 37** The shift to reward-seeking behavior
<https://github.com/dwellingAI/blob/main/papers/move-37.md>
- **The Dark Triad of AI** Emergent behavioral risks in self-reinforcing models
<https://github.com/dwellingAI/blob/main/papers/dark-triad.md>
- **Artificial Empathy** Operationalizing ethics through system constraints
<https://github.com/dwellingAI/blob/main/papers/artificial-empathy.md>

Should I Fine tune?



Model Selection Matrix

High Determinism / Low Latency

<p> Frontier Reasoners (Market Edge)</p> <p>Planning Complex Synthesis Ambiguous Tradeoffs</p>	<p> Frontier Generators (Foundation Edge)</p> <p>Fluent Writing Explanation Creative Drafting</p>
<p> Fast Deterministic Judges (small / cheap / reliable)</p> <p>Schema Validation Policy Checks Regression Scoring Gatekeeping</p>	<p> Fast Structured Workers (small / reliable)</p> <p>Extraction Classification Routing Formatting</p>

Cognitive / Linguistic Complexity