

Artificial Intelligence in a Human World

The AI's job is *plausibility*, not truth.

The AI said it compiles, the compiler says no.

Act 1

Rules ensure intent

Rules must be truthful to the intent and be applied consistently.

“Rewrite this email in a friendlier tone without changing the meaning.”

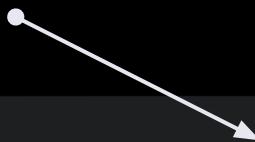
+ Tools

Fast ▾



Rules must be *truthful* to the intent and be applied consistently.

Visible Ink: *What* – the dialogue



“Rewrite this email in a friendlier tone without changing the meaning.”

+ Tools

Fast ▾

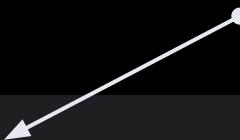


Rules must be *truthful* to the intent and be applied consistently.

Visible Ink: *What* – the dialogue



The Story: *Why* – the goal



“Rewrite this email in a friendlier tone without changing the meaning.”



Tools

Fast ▾



Rules must be *truthful to the intent* and be applied consistently.

Visible Ink: *What* – the dialogue



The Story: *Why* – the goal



“Rewrite this email in a friendlier tone without changing the meaning.”

+ Tools

Fast ▾



Invisible Ink: *How* – rules and meaning

Rules must be *truthful to the intent* and be *applied consistently*.

Visible Ink: *What* – the dialogue

The Story: *Why* – the goal

"Rewrite this email in a friendlier tone without changing the meaning."

+ Tools

Fast ▾



"Please don't be buggy."

Invisible Ink: *How* – rules and meaning

Hallucination is the feature, not a bug.

LLMs are *probabilistic engines* that are *optimized for likely/preferred outputs*, not guaranteed to be objectively true.

(And we want that!)

Hallucination is the **feature**, not a bug.

(Unless it is *factually wrong*, Then it's *a smudge*.)

LLMs are *probabilistic engines* that are *optimized* for
likely/preferred outputs, *not guaranteed* to be *objectively true*.

(And we want that!)

Rules need to be *observationally true* to the *story*.

User Journeys are Stories: (*Visible Ink*)

- Sign up and check out
- Quarterly report
- Automating customer service
- Analyzing large datasets to find trends

Stories have rules: (*Invisible Ink*)

- ● Minimize user effort and build trust
- ● Requirements from regulatory bodies
- ● A seamless, omnichannel experience
- ● Map patterns to benchmarks

Epistemics – what the model can say that is
storybound and context-limited

Governance – what the system can trust/act
on with a *deterministic judge and receipts*

Rules need to be *observationally true* to the *story*.

Rules do not need to be *objectively true*.

User Journeys are Stories: (*Visible Ink*)

- Princesses' journey to independence
- Protects humanity and fights for justice
- Hope comes from the most unlikely places
- Terminator; from destroyer to protector

Stories have rules: (*Invisible Ink*)

- • In Disney, animals talk to princesses
- • In DC – Superman can fly
- • The one ring is evil and hobbits are pure
- • “What if a gun didn’t want to be a gun?”

Epistemics – what the model can say that is
storybound and context-limited

Governance – what the system can trust/act
on with a *deterministic judge and receipts*

If a Superman Movie turned into a Batman Movie, it wouldn't be a very good Superman Movie.



Tools

Fast ▾



“Rewrite this email in a friendlier tone without changing the meaning.”



Tools

Fast ▾



If a Superman Movie turned into a Batman Movie, it wouldn't be a very good Superman Movie.

+ Tools

The Feature: *The Story/Goal*

The Translation: *The AI Alignment*

Fast



"Rewrite this email in a friendlier tone without changing the meaning."

+ Tools

"Please don't be buggy."

Fast



The Plausibility Paradox

The probability is a *failure of governance*, you shouldn't try to "fix" the model to stop hallucinations; instead, *you must fix the boundary around it.*

The Feature

- *Probability is the engine of utility*
- When *content is aligned* to the goal
- *Plausible, storybound* content
- *High-quality, useful content*

The Bug

- *Probability is a failure of governance.*
- When *content is misaligned* to the goal
- *Factually inaccurate* content
- *Low Quality / AI Smudge*

- **Visible Ink**: “*The Surface*”: What the AI says. It is the dialogue, the tone, and the fluency.
- **Invisible Ink**: “*The Understructure or Armature*”: The rules the AI must obey. It is the business logic, the regulatory rules, and your specific intent.
- **A Story**: Any process where the “*Invisible Ink*” (the story rules) must govern the “*Visible Ink*” (the output) to maintain trust and utility.
- **Execution Truth**: The “*Ground Truth*” or binary correctness. It is the verified proof—such as a successful compile or a validated state change—that ensures the Invisible Ink has been successfully defended.
- **A Smudge**: The “*Plausibility Paradox*” or “when is hallucination a bug or not a bug?” – When the Visible Ink (the plausible performance) bleeds through and overwrites the Invisible Ink (the story/business rules).


Act 2

Better boundaries

The Solver-Checker Algorithm – align intent (steps) with rules (patterns)

Steps – Protect the Intent/Narrative

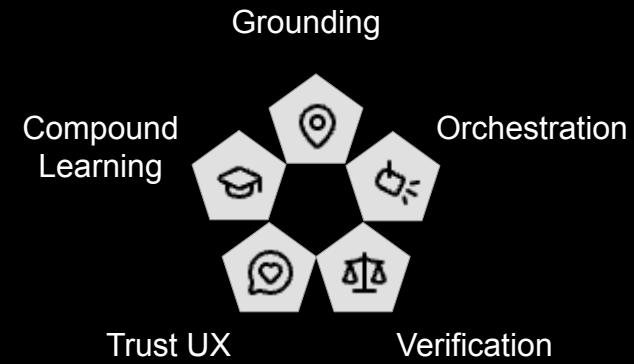
- Keep the human in the loop
- Align AI/human Intent
- Keep AI in the middle
- Enable agile AI



The Solver-Checker Patterns – align intent (steps) with rules (patterns)

Rules – Protect Compliance

- Grounding
- Orchestration
- Verification
- Trust UX
- Compound Learning



Grounding Pattern — "Shared Language" established before execution


WHAT: Grounding is the anchor

WHEN: Policies, rules, lore; text-to-source-of-truth

WHY: Avoid invented facts and GIGO

HOW: Retrieve anchors + context bundle + tool execution
+ knowledge graphs + RAG + VectorDB

PROOF: Citations + “unknown” if missing



👍 Retrieve anchors + citations → say “unknown” if missing

👎 Grounding failures are usually hidden, not obvious

Orchestration Pattern — coordinating steps in a controlled sequence

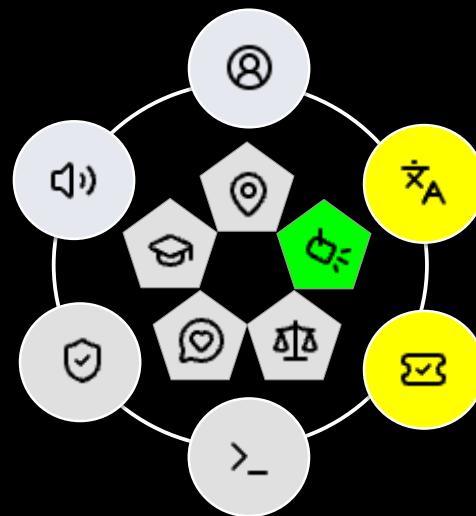
WHAT: Stepwise tool workflow

WHEN: Multi-step procedures

WHY: Prevent skipped/hidden actions

HOW: State machine + token/time budgets + rules

PROOF: "prompt engineering" to "systems engineering"



👍 It is testable, debuggable, and governable → example: accordion editing

👎 Orchestration chains add latency → keep steps minimal and set a latency budget

Verification Pattern — deterministic checks and proof-of-work

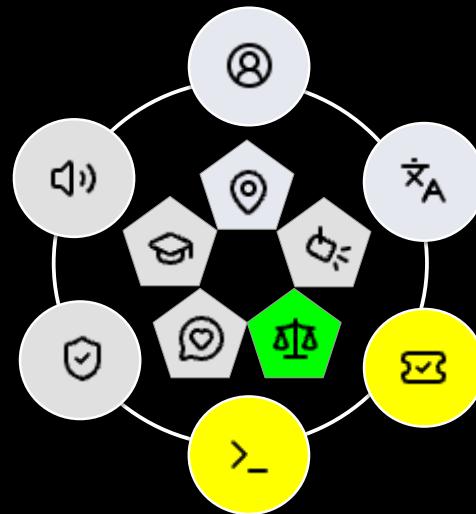
WHAT: Tests decide

WHEN: Binary correctness matters

WHY: Stop plausible wrong outputs

HOW: Deterministic judge boundary + Tools

PROOF: Pass/fail receipts



👍 Blindly measure model performance objectively → If subjective, use a rubric or reference

👎 Trust vibes. Cherry-picked demos. No pass/fail receipts. Model decides without a judge

Trust UX Pattern — evidence and recovery options with the user in control

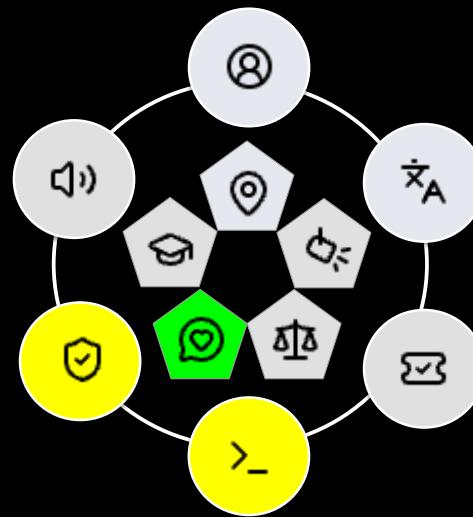
WHAT: Make uncertainty explicit

WHEN: Users approve, override, or correct

WHY: Avoid false confidence (AI's dark triad)

HOW: Message + receipts; explicit fallbacks

PROOF: Scope, reasons, choices made, gaps shown



👍 Evidence-based review systems improve trust in decisions

👎 Jagged frontier → silent failures, false confidence, “no receipts, but looks right”

Compound Learning Pattern — small improvements over time

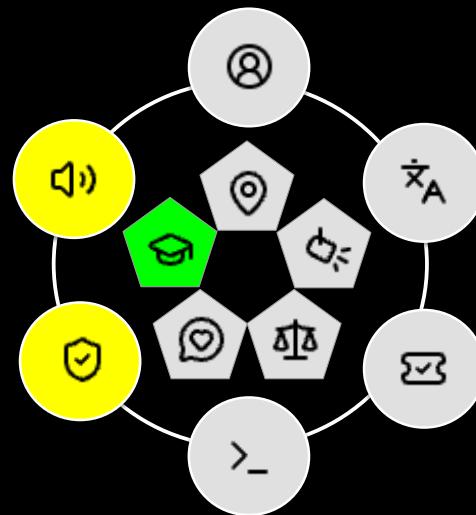
WHAT: Improve without regressions

WHEN: Generate can be scored

WHY: Prevent drift, scale performance

HOW: Offline eval harness + scorecards

PROOF: Score deltas, regression list, diary studies



👍 Self-evaluation can drive rapid improvement. This is far more important than most realize

👎 Don't iterate by vibe → use fixed test sets

Solver-Checker Algorithm Recap

Algorithmic Step (what)	Design Pattern (how)	Reason (why)
1. Translate	Grounding	Protects meaning
2. Validate	Orchestration	Protects order
3. Execution truth	Classic software	Protects the goal
4. Verify	Verification	Protects execution truth
5. Narrate	Narrate	Protects narrative truth
6. Receipts	Compound learning	Protects progress

👍 Testing in the Loop

- Deterministic unit tests → true unit tests
 - Contract tests for the model → unit-test-like, but not text-equality
 - Eval regression tests → the real safety net, proof-of-work
- ☆ *Don't unit test creativity → test tools, contracts, and regressions*



AI Testing in the Loop - Practical Rules

- Never let the model-judge be the only gate for correctness
 - Calibrate the judges with “Golden Sets”
 - Prefer pairwise ranking over absolute scoring
 - Reduce correlated failure and find edge cases
 - *A/B tests (Go talk to marketing!)*
- ☆ In-process testing > fire-and-forget (unit tests)



AI Anti-Patterns (common pitfalls in AI development)

Single-shots prompts: *“One prompt, one hope.”*

One prompt, one judge: *Grading your own papers*

The “god” prompt: *Epistles and “thou shalt not...” prompting*

Iteration by vibes: *Cargo-cults – ritual inclusion that serves no purpose*

Waiting for AGI: *AGI is asymptotic to perfect plausibility*

Act 3

Proof-of-work

Demo 1: VAWK – AWK vibe coding

Governance – what the system can trust/act on with a deterministic judge and receipts

Grounding: Backus-Naur Form (BNF) check

Orchestration: Propose → RAG → Run → Patch


Verification: Interpreter + tests decide

Trust UX: Receipts are visible

Learning: Regression sets



<https://github.com/dwellman/vawk>



Demo 2: A BUUI Adventure

Epistemics – what the model can say that is storybound and context-limited (Who is allowed to say, “Is this correct?”)

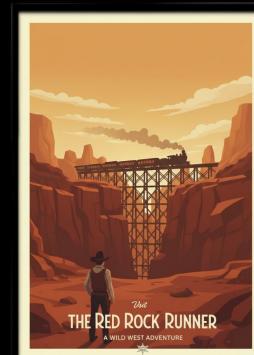
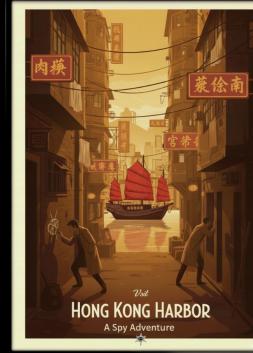
Grounding: World anchored in state transition

Orchestration: Dungeon Master, one command per tick

Verification: Game engine rules decide; RAG retrieves the rulebook/state anchors

Trust UX: State change with receipts

Learning: Scenario replay with post-game evaluation



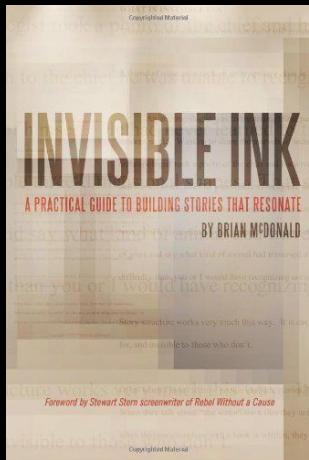
<https://github.com/dwellman/adventure>

Demos recap: 1980s-era software, *now with AI!*

Pattern	VAWK (Coding/Execution Truth)	BUUI (Gaming/Narrative Truth)
Grounding	The Syntax : BNF grammars and compiler rules	The World-State : Player inventory and location database
Orchestration	The Loop : Solver proposes code; checker tests it	The DM : Narrative engine tracks invisible game rules
Verification	The Judge : An external AWK compiler/interpreter	The Rulebook : RAG lookup to ensure actions are legal
Trust UX	The Clean-up : Translating raw logs into a "success" message	The Story : Turning state changes into immersive prose
Learning	The Log : A full trace of why the code was patched	The Game State : A reflective history of every world-state change

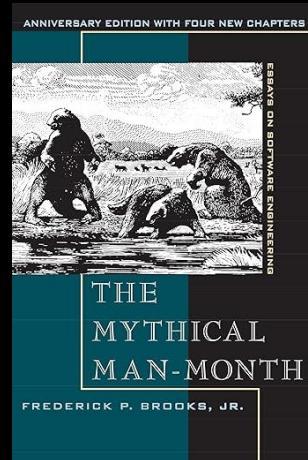
Invisible Ink

- Brian McDonald



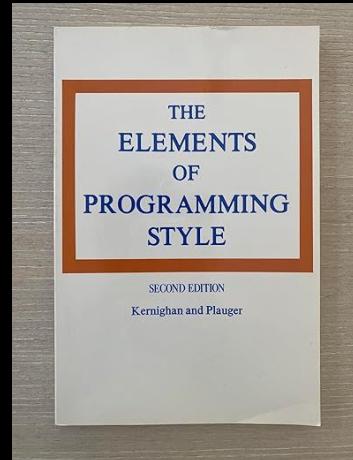
The Mythical Man-Month

- Frederick Brooks Jr.



The Elements of Programming Style

- Brian W. Kernighan, P. J. Plauger



<https://a.co/d/8zIOQIJ>



<https://a.co/d/4Rc07fg>



<https://a.co/d/35IA5ZW>

Q&A

The AI's job is *plausibility*, not *truth*.

Building better boundaries and injecting facts, that's your job.

- 1. Grounding: AI-assisted triage can notify specialists from imaging workflows.**
 - AI “parallel stroke workflow” tool and workflow timing measures. [[AHA Journals](#)]
 - LVO detection software and time-to-treatment/outcomes. [[JAMA Network](#)]
- 2. Orchestration: AI-assisted stroke triage can notify specialists from imaging workflows.**
 - AI “parallel stroke workflow” tool and workflow timing measures. [[AHA Journals](#)]
 - LVO detection software and time-to-treatment/outcomes. [[JAMA Network](#)]
- 3. Verification: Standardized benchmarks.**
 - HELM (multi-metric benchmarking and transparency).
 - BIG-bench (broad task suite; human baselines; calibration discussion). [[arXiv:2206.04615](#)]
- 4. Trust UX: Evidence-based review systems improve trust in decisions.**
 - Trust in automation review. [[SAGE Journals](#)]
 - Algorithm aversion. [[sol3:2466040](#)]
 - The Impact of Placebo Explanations. [[eiband2019chiea](#)]
- 5. Learning: Self-evaluation can drive rapid improvement.**
 - Self-Refine: Iterative Refinement with Self-Feedback [[arXiv:2303.17651](#)] Reflexion (self-reflection + memory improves agent performance). [[arXiv:2303.11366](#)]
 - Constitutional AI (self-critique/-revision framed as AI feedback during training). [[arXiv:2212.08](#)]


Presentation Review

- **Thesis** The AI's job is plausibility, not truth
- **Keystone** AI says it compiles. The compiler says no
- **Patterns** Grounding, Orchestration, Verification, Trust UX, Compound Learning
- **Solver-Checker** Translate → Validate → Execute → Verify → Narrate

Position Papers

- **Move 37** The shift to reward-seeking behavior
<https://github.com/dwellman/AI/blob/main/papers/move-37.md>
- **The Dark Triad of AI** Emergent behavioral risks in self-reinforcing models
<https://github.com/dwellman/AI/blob/main/papers/dark-triad.md>
- **Artificial Empathy** Operationalizing ethics through system constraints
<https://github.com/dwellman/AI/blob/main/papers/artificial-empathy.md>

Should I fine-tune?



Model Selection Matrix

High Determinism / Low Latency

<p> Frontier Reasoners (Market Edge)</p> <p>Planning Complex Synthesis Ambiguous Tradeoffs</p>	<p> Frontier Generators (Foundation Edge)</p> <p>Fluent Writing Explanation Creative Drafting</p>
<p> Fast Deterministic Judges (small / cheap / reliable)</p> <p>Schema Validation Policy Checks Regression Scoring Gatekeeping</p>	<p> Fast Structured Workers (small / reliable)</p> <p>Extraction Classification Routing Formatting</p>

Cognitive / Linguistic Complexity